

Certainly! Let's dive deeper into YARN (Yet Another Resource Negotiator), focusing on its architecture, components, and a technical example.

YARN Architecture Overview

YARN is designed to provide a more scalable and flexible way of managing the computing resources in the Hadoop cluster. It separates the functionalities of resource management and job scheduling/monitoring into separate daemons: the ResourceManager (RM), the NodeManager (NM), and the ApplicationMaster (AM).

1. ResourceManager (RM):

- The RM is the master daemon of YARN and is responsible for allocating resources to various running applications.
- It has two main components:
 - **Scheduler:** Responsible for allocating resources to various running applications, based on constraints such as capacity, queues, etc.
 - **ApplicationManager:** Manages the lifecycle of applications and is responsible for accepting job submissions, negotiating the first container for executing the application-specific ApplicationMaster, and provides the service for restarting the AM container on failure.

2. NodeManager (NM):

- A per-node slave daemon, it is responsible for monitoring resource usage (CPU, memory, disk, network) and reporting the same to the ResourceManager.
- It is also responsible for managing the application's running containers, ensuring they are executed with the resource constraints (memory/CPU) specified.

3. ApplicationMaster (AM):

- An instance of the ApplicationMaster is created for each application. The AM is responsible for negotiating appropriate resource containers from the Scheduler, tracking their status, and monitoring progress.

Example: Running a MapReduce Job on YARN

Prerequisites:

- A Hadoop cluster with YARN configured.
- Sample data and MapReduce job (e.g., word count).

Steps:

1. Prepare the Environment:

- Place your input data in HDFS:
`hdfs dfs -put localfile.txt /user/hadoop/input`

2. Write a MapReduce Program:

- For simplicity, let's assume a basic word count program in Java.
3. **Compile and Package the Application:**
 - Use Maven or another build tool to compile your Java code and package it into a JAR.
 4. **Submit the Job:**
 - Use the `yarn` command to submit the job to the cluster.
`yarn jar your-application.jar com.example.WordCount /user/hadoop/input /user/hadoop/output`
 5. **Monitoring the Job:**
 - You can monitor the job's progress through the YARN Resource Manager UI.

Example Code (WordCount.java):

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import java.io.IOException;
import java.util.StringTokenizer;

public class WordCount {

    public static class TokenizerMapper extends Mapper<Object, Text, Text, IntWritable>{
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }

    public static class IntSumReducer extends Reducer<Text,IntWritable,Text,IntWritable> {
        private IntWritable result = new IntWritable();

        public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException,
            int sum = 0;
    }
}
```

```

        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```